

SHEHAB MOHAMED IBRAHIM

Cairo, Egypt • +20 1153544745 • shehabsdm@gmail.com • [LinkedIn](#) • [GitHub](#)

PROFESSIONAL SUMMARY

AI Engineer, Machine Learning Engineer, and Data Scientist with a proven track record delivering production-grade systems for enterprise clients across the US, Canada, and Singapore. Specializes in Large Language Models (LLM), Retrieval-Augmented Generation (RAG), Applied AI, and Multi-Agent Orchestration — with deep expertise in Natural Language Processing (NLP), Deep Learning, and Computer Vision. Delivered 0.14 ms P95 API latency and 60% faster model deployment in production environments on GCP and AWS.

PROFESSIONAL EXPERIENCE

- AI & Machine Learning Engineer** • *Niibu Inc* • *New York, USA* — *Remote (Full-time)*

Jan 2026 – Present

 - Architected the "Neural Concierge" Multi-Agent LLM Orchestration system on GCP using Docker, processing 290K+ records to deliver real-time AI-driven personalized matching via XGBoost (rank:pairwise) and Deep Learning
 - Engineered the "DNA Affinity" ML ranking engine (40% Loyalty / 25% Frequency / 20% Behavior / 15% Context), achieving 0.14 ms P95 API latency by implementing gRPC, Feast Feature Store, and Redis caching
 - Hardened platform security and ensured HIPAA-aligned healthtech compliance by implementing WireGuard tunneling, RS256 JWT authentication, and Tenant Isolation — fully documented via Swagger (OpenAPI)
- MLOps & AI Engineer** • *MacroTrend Analytics* • *Canada* — *Remote (Part-time)*

Aug 2025 – Dec 2025

 - Accelerated model time-to-production by 60% for 500+ users by deploying Deep Learning forecasting models on Kubernetes with robust Continuous Training (CT) pipelines
 - Optimized real-time inference to <50 ms at 1,000+ req/sec by integrating ONNX/TensorRT and FastAPI, reducing serving costs and improving throughput
 - Established 100% model drift auditability across all production ML systems by building a Prometheus, Grafana, and MLflow observability stack
- AI Solutions Engineer** • *Flowmingo AI* • *Singapore* — *Remote (Part-time)*

May 2025 – Jul 2025

 - Reduced inference latency by 40% for enterprise B2B clients by fine-tuning and deploying LLM chatbots using LangChain on AWS, applying Applied AI optimization techniques including quantization and prompt engineering
 - Eliminated 50% of manual recruitment screening and boosted hiring efficiency by 35% by designing a Multi-Agent automation architecture for end-to-end candidate evaluation
- AI Agent Engineer Intern** • *NeuralSeek* • *United States* — *Remote*

Jan 2025 – Mar 2025

 - Reduced LLM hallucination rate by 35% and achieved 100% source-grounded responses by deploying FAISS Vector Database-backed RAG systems with Human-in-the-Loop (HITL) validation integrated into multi-LLM Llama-3 workflows
- Data Science Intern** • *PM Accelerator* • *United States* — *Remote*

Aug 2024 – Oct 2024

 - Enforced full HIPAA compliance across 1.5 TB of sensitive healthcare and medical imaging datasets by building end-to-end RAG and NLP pipelines with SHA-256 data anonymization protocols — leveraging Biomedical Engineering domain expertise

TECHNICAL SKILLS

- Generative AI & LLMs:** GPT-4, Llama-3, RAG Architecture, Multi-Agent Orchestration, Fine-Tuning (QLoRA, PEFT, SFT, DPO), LangChain, OpenAI API, Hugging Face, Prompt Engineering
- Machine Learning & Deep Learning:** PyTorch, TensorFlow, Scikit-learn, XGBoost, Transformers, FAISS Vector Database, ONNX/TensorRT, Time-Series Forecasting, LSTM
- Natural Language Processing (NLP):** Semantic Search, Embeddings, Text Classification, Named Entity Recognition (NER), Chatbot Development, Hallucination Reduction, HITL Validation
- Computer Vision:** OpenCV, YOLO, Image Classification, Object Detection, Medical Imaging, Medical Image Analysis, Statistical Modeling
- MLOps & Cloud:** GCP, AWS (EKS/S3/EC2), Docker, Kubernetes, Terraform, MLflow, DVC, Prometheus, Grafana, Feast Feature Store, CI/CD, Continuous Training Pipelines
- Backend & APIs:** Python, FastAPI, gRPC, RESTful APIs, Redis, SQL/NoSQL, Pandas, NumPy, Swagger/OpenAPI
- Security & Compliance:** HIPAA Compliance, RS256 JWT, WireGuard Tunneling, Tenant Isolation, SHA-256 Anonymization

CERTIFICATIONS & COURSES

Hugging Face NLP Course — Hugging Face | DeepLearning.AI — LangChain for LLM Application Development — DeepLearning.AI | Machine Learning Specialization — Coursera / Andrew Ng

EDUCATION

- B.Sc. Biomedical Engineering — AI & Machine Learning Focus** • *Cairo University, Egypt*

June 2025
- Relevant Coursework:** Machine Learning, Deep Learning, Natural Language Processing (NLP), Computer Vision, Neural Networks, Statistical Modeling, Data Structures & Algorithms